
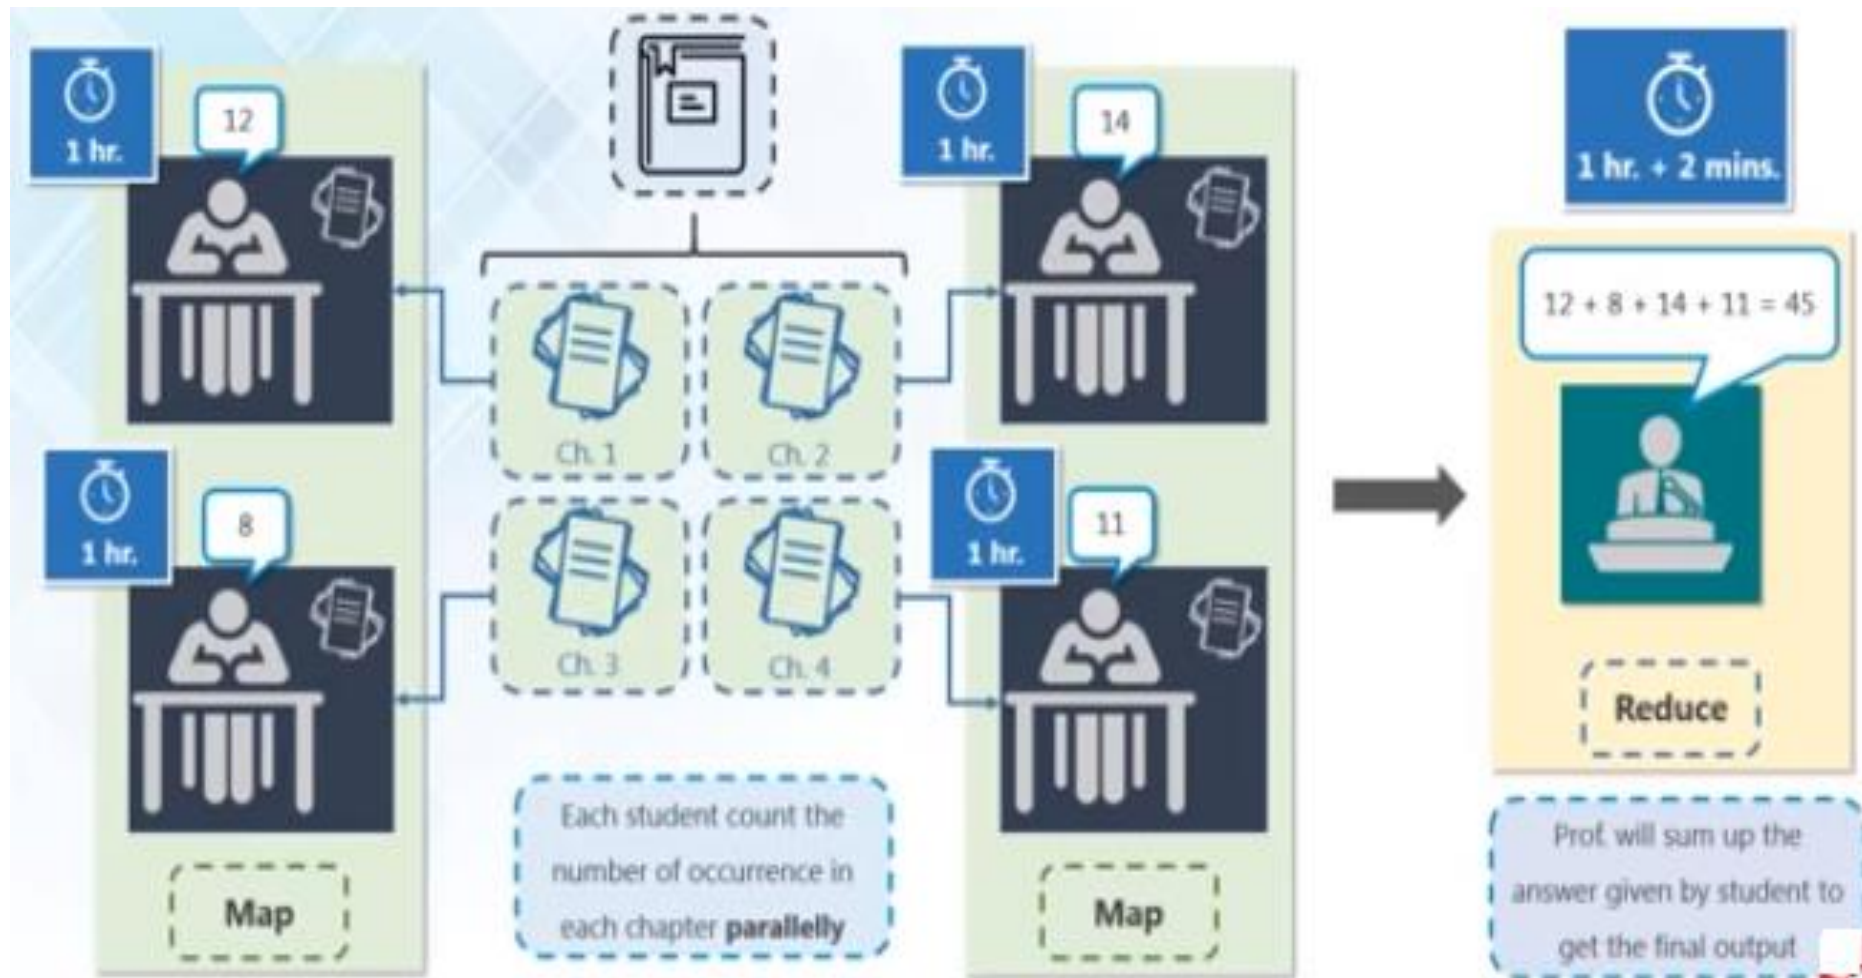


MapReduce

Story of MapReduce




Story of MapReduce

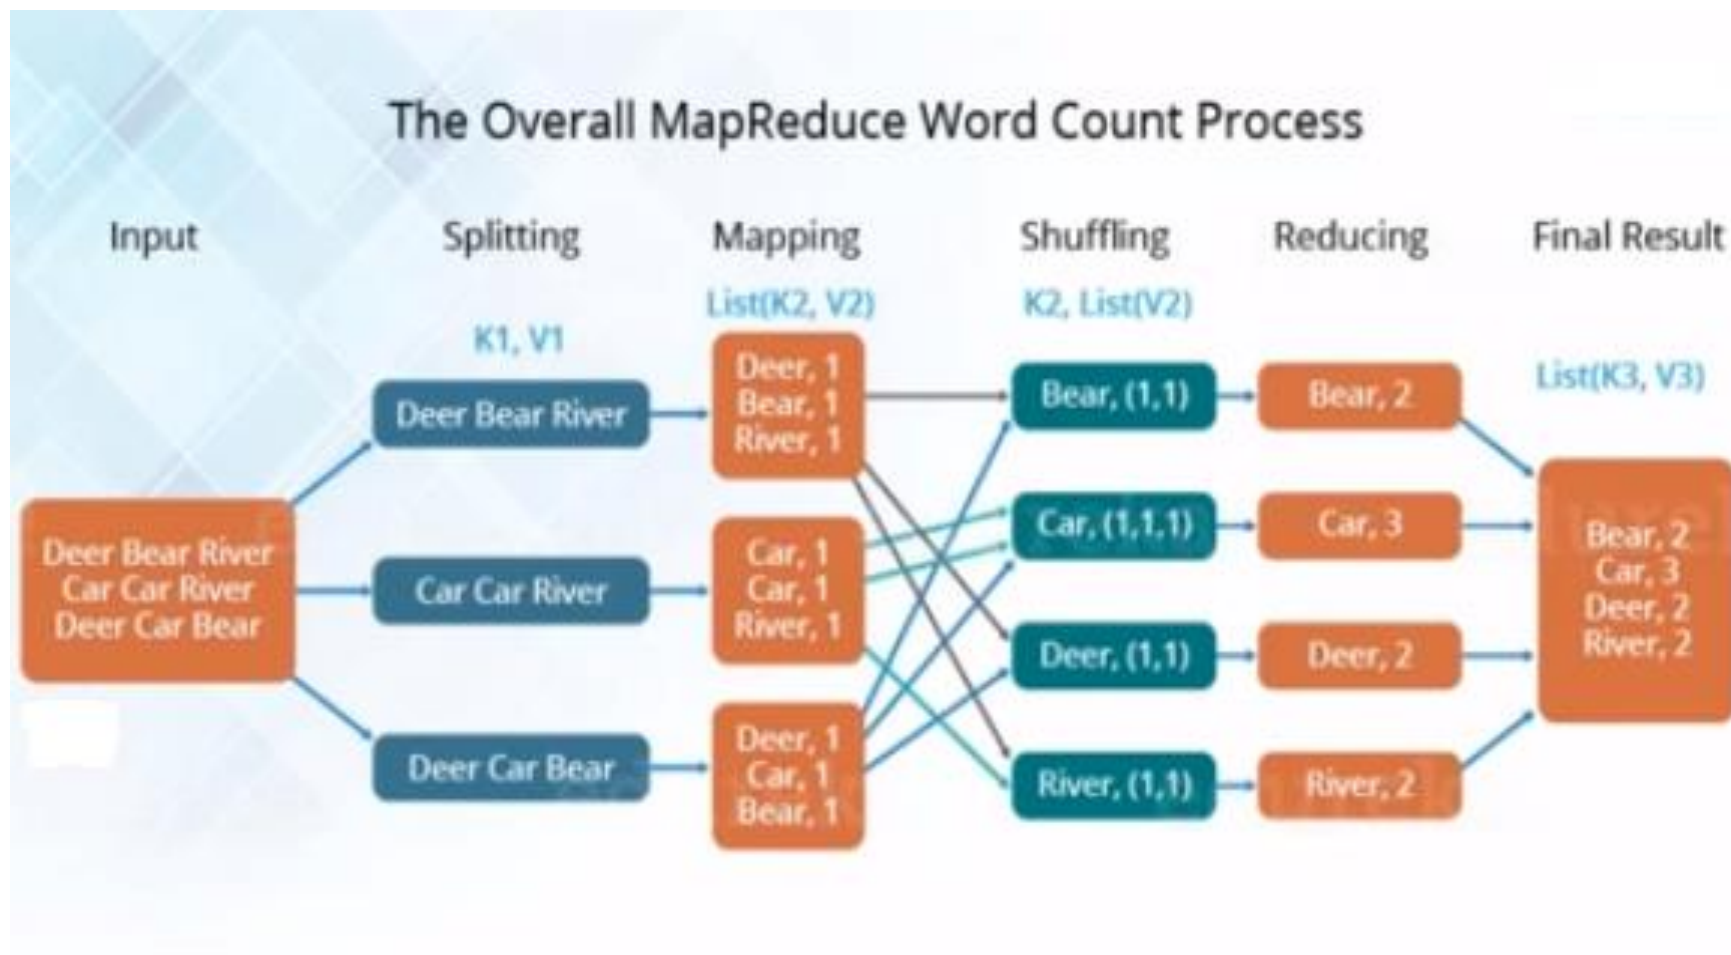


What is MapReduce?

MapReduce is a programming framework that allows us to perform distributed and parallel processing on large data sets in a distributed environment



MapReduce Word Count Example



MapReduce Wordcount Example

Three Major Parts of MapReduce Program:

1

Mapper Code:

You write the mapper logic over here i.e. how map task will process the data to produce the key-value pair to be aggregated

2

Reducer Code:

You write reducer logic here which combines the intermediate key-value pair generated by Mapper to give the final aggregated output

3

Driver Code

You specify all the job configurations over here like job name, Input path, output path, etc.

Packages and classes

```
import java.io.IOException;  
import java.util.*;
```

```
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.conf.*;  
import org.apache.hadoop.io.*;
```

**All these packages are present in
hadoop-common.jar**


```
import org.apache.hadoop.mapreduce.*;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;  
import  
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
import  
org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

**All these
packages are
present in
hadoop-mapreduce-
client-core.jar**


Mapper Class

```
public static class Map extends  
Mapper<LongWritable, Text, Text, IntWritable> {
```

**Name of the Mapper Class which
inherits Super Class Mapper**

A blue arrow points from the text 'Name of the Mapper Class which inherits Super Class Mapper' to the 'Map' class name in the code snippet above.

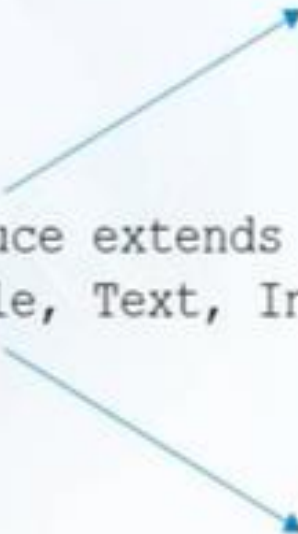
**Mapper Class takes 4 Arguments i.e.
Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT>**

A blue arrow points from the text 'Mapper Class takes 4 Arguments i.e. Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT>' to the generic type arguments in the code snippet above.

Reducer Class

```
public static class Reduce extends  
Reducer<Text, IntWritable, Text, IntWritable> {
```

**Name of the Reducer Class which
inherits Super Class Reducer**



**Reducer Class takes 4 Arguments i.e.
Reducer <KEYIN, VALUEIN, KEYOUT, VALUEOUT>**

Byte Offset

- Byte offset is the number of character that exists counting from the beginning of a line.
- Byte offset is represented hexadecimal.
- E.g. this line “what is byte offset” will have a byte offset of 19. This is used as key value in hadoop
- Key = byte offset, value = line

Mapper Code



```
public static class Map extends Mapper<LongWritable,Text,Text,IntWritable> {  
  
    public void map(LongWritable key, Text value, Context context) throws IOException,InterruptedException {  
  
        String line = value.toString();  
        StringTokenizer tokenizer = new StringTokenizer(line);  
        while (tokenizer.hasMoreTokens()) {  
            value.set(tokenizer.nextToken());  
            context.write(value, new IntWritable(1));  
        }  
    }  
}
```


Mapper Input:

- The key is nothing but the offset of each line in the text file: *LongWritable*
- The value is each individual: *Text*

Mapper Output:

- The key is the tokenized words: *Text*
- We have the hardcoded value in our case which is 1: *IntWritable*
- Example - Dear 1, Bear 1, etc.

Reducer Code



Reducer Input:

- Keys are unique words which have been generated after the sorting and shuffling phase: Text
- The value is a list of integers corresponding to each key: IntWritable
- Example: Bear, [1, 1], etc.

Reducer Output:

- The key is all the unique words present in the input text file: Text
- The value is the number of occurrences of each of the unique words: IntWritable
- Example: Bear, 2; Car, 3, etc. .

Driver Code

In the driver class, we set the configuration of our MapReduce job to run in Hadoop

```
Configuration conf= new Configuration();
Job job = new Job(conf,"My Word Count Program");
job.setJarByClass(WordCount.class);
job.setMapperClass(Map.class);
job.setReducerClass(Reduce.class);
job.setOutputKeyClass(Text.class);

job.setOutputValueClass(IntWritable.class);
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);
Path outputPath = new Path(args[1]);

//Configuring the input/output path from the filesystem into the job
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

- Specify the name of the job , the data type of input/output of the mapper and reducer
- Specify the names of the mapper and reducer classes.
- Path of the input and output folder
- The method `setInputFormatClass ()` is used for specifying the unit of work for mapper
- `Main()` method is the entry point for the driver

MapReduce Word Count Example

What is Task in Map Reduce?

- A task in MapReduce is an execution of a Mapper or a Reducer on a slice of data.
- It is also called Task-In-Progress (TIP).
- It means processing of data is in progress either on mapper or reducer.

What is Task Attempt?

- Task Attempt is a particular instance of an attempt to execute a task on a node.
- There is a possibility that anytime any machine can go down.
- For example, while processing data if any node goes down, framework reschedules the task to some other node. This rescheduling of the task cannot be infinite.
- There is an upper limit for that as well. The default value of task attempt is 4. If a task (Mapper or reducer) fails 4 times, then the job is considered as a failed job.
- For high priority job or huge job, the value of this task attempt can also be increased.

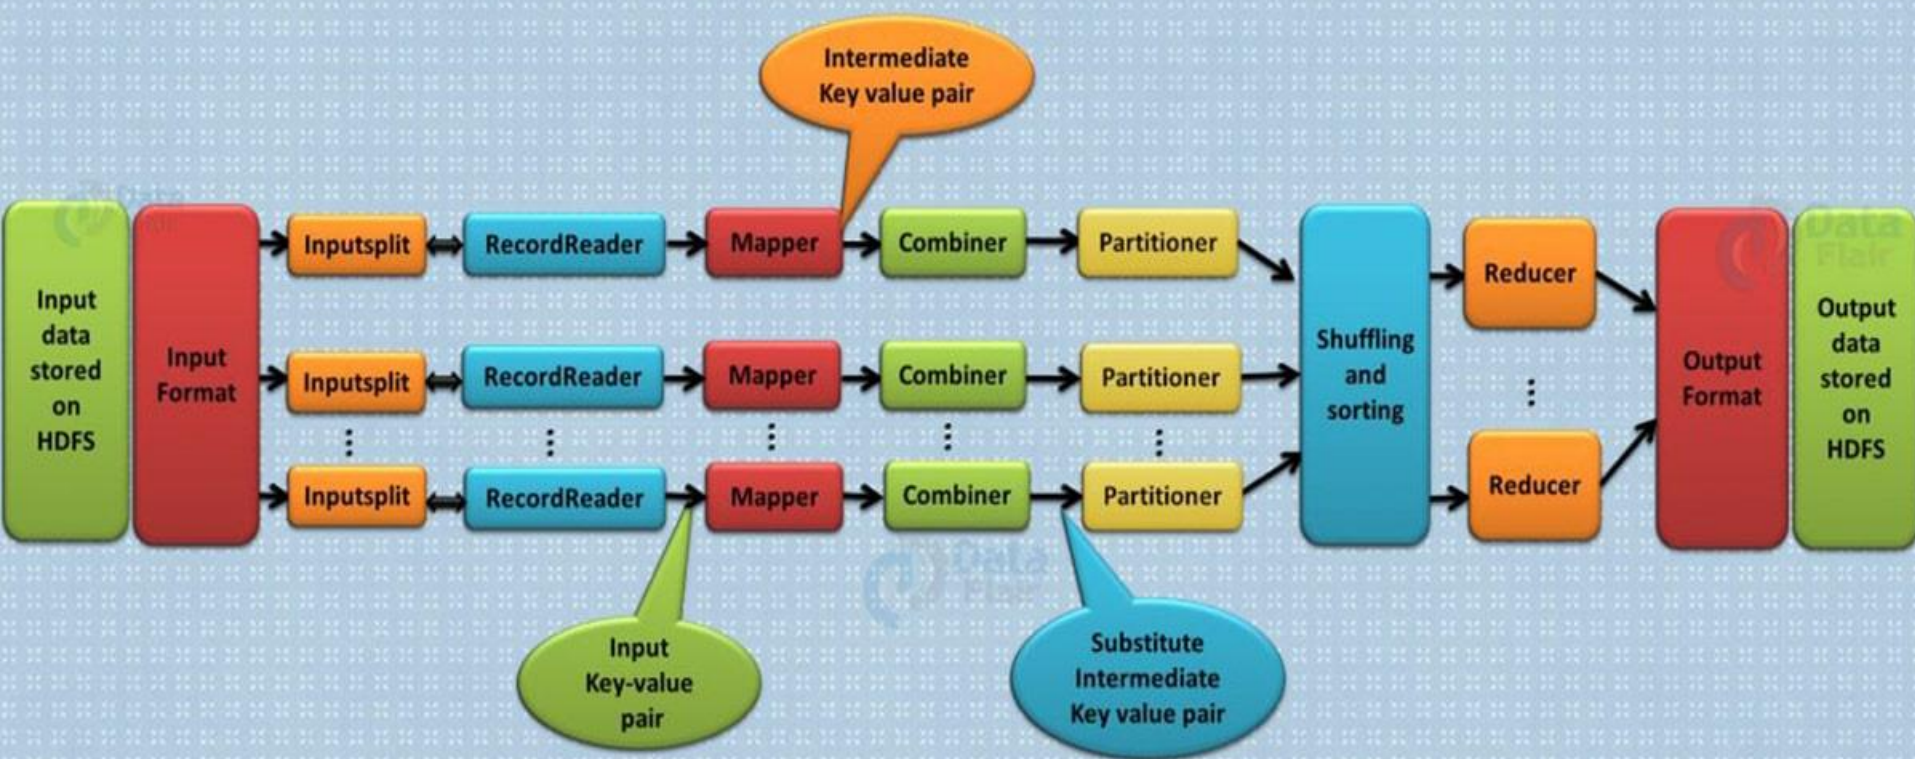
Failed Task v/s Killed Task

- A failed task attempt is a task attempt that completed, but with an unexpected status value.
- A killed task attempt is a duplicate copy of a task attempt that was started as part of speculative execution.
- Hadoop uses "speculative execution." The same task may be started on multiple boxes. The first one to finish wins, and the other copies are killed.
- Failed tasks are tasks that error out.
- There are a few reasons Hadoop can kill tasks by its own decisions:
 - Task does not report progress during timeout (default is 10 minutes)
 - FairScheduler or CapacityScheduler needs the slot for some other pool (FairScheduler) or queue (CapacityScheduler).
 - Speculative execution causes results of task not to be needed since it has completed on other place.

MapReduce Phases

- Mapper
 - RecordReader
 - Map
 - Combiner
 - Partitioner
- Reducer
 - Shuffle
 - Sort
 - Reduce
 - Output Format

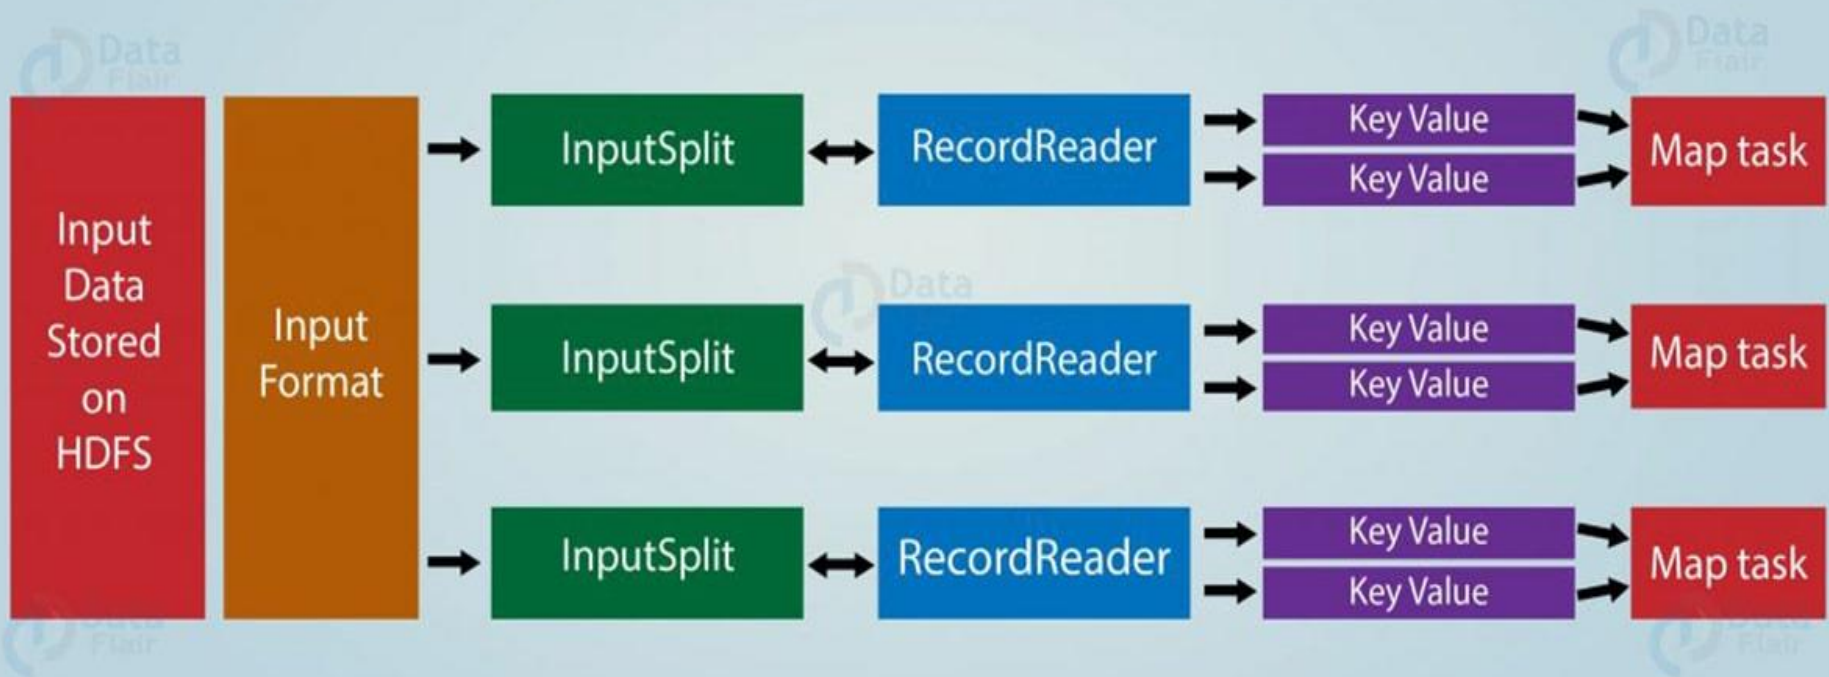
MapReduce Job Execution Flow



Key-Value Pairing

- The output types of the Map should match the input types of the Reduce as shown below:
 - **Map:** $(K1, V1) \rightarrow \text{list}(K2, V2)$
 - **Reduce:** $\{(K2, \text{list}(V2))\} \rightarrow \text{list}(K3, V3)$

Key Value Pairing in Hadoop MapReduce



Key – Value Pairing

- Generation of a key-value pair in Hadoop depends on the data set and the required output.
- In general, the key-value pair is specified in 4 places:
 - Map input,
 - Map output,
 - Reduce input,
 - Reduce output

Key – Value Pairing

- **a. Map Input**

- Map-input by default will take the line offset as the key and the content of the line will be the value as Text. By using custom InputFormat we can modify them.

- **b. Map Output**

- Map basic responsibility is to filter the data and provide the environment for grouping of data based on the key.
- **Key** – It will be the field/ text/ object on which the data has to be grouped and aggregated on the reducer side.
- **Value** – It will be the field/ text/ object which is to be handled by each individual reduce method.

- **c. Reduce Input**

- The output of Map is the input for reduce, so it is same as Map-Output.

- **d. Reduce Output**

- It depends on the required output.

Input Files

- The data for a MapReduce task is stored in **input files**, and input files typically lives in **HDFS**.
- The format of these files is arbitrary, while line-based log files and binary format can also be used.

InputFormat

- **InputFormat** defines how input files are split and read.
 - e.g. Line by line
- It selects the files or other objects that are used for input.
- InputFormat creates InputSplit.

Functionality of InputFormat

- The files or other objects that should be used for input is selected by the InputFormat.
- InputFormat defines the Data splits, which defines both the size of individual Map Tasks and its potential execution server.
- InputFormat defines the RecordReader, which is responsible for reading actual records from the input files.

Types of InputFormat in MapReduce

FileInputFormat

TextInputFormat

KeyValueTextInputFormat

SequenceFileInputFormat

SequenceFileAsTextInputFormat

NLineInputFormat

SequenceFileAsBinaryInputFormat

DBInputFormat



InputSplits

- Created by InputFormat
- Logically represents the data which will be processed by an individual **Mapper**.
- One map task is created for each split
- So, no. of map tasks = no. of InputSplits.
- The split is divided into records and each record will be processed by the mapper.
- Length of InputSplit is measured in bytes

InputSplit vs Block

- **Block** – The default size of the HDFS block is 128 MB which we can configure as per our requirement. All blocks of the file are of the same size except the last block, which can be of same size or smaller. The files are split into 128 MB blocks and then stored into Hadoop FileSystem.
- **InputSplit** – By default, split size is approximately equal to block size. InputSplit is user defined and the user can control split size based on the size of data in MapReduce program.

Data Representation

- **Block** – It is the physical representation of data. It contains a minimum amount of data that can be read or write.
- **InputSplit** – It is the logical representation of data present in the block. It is used during data processing in MapReduce program or other processing techniques. InputSplit doesn't contain actual data, but a reference to the data.

InputSplit vs Block

Example.txt

130 MB

File split into
2 blocks

Block 1

128 MB

Block
2

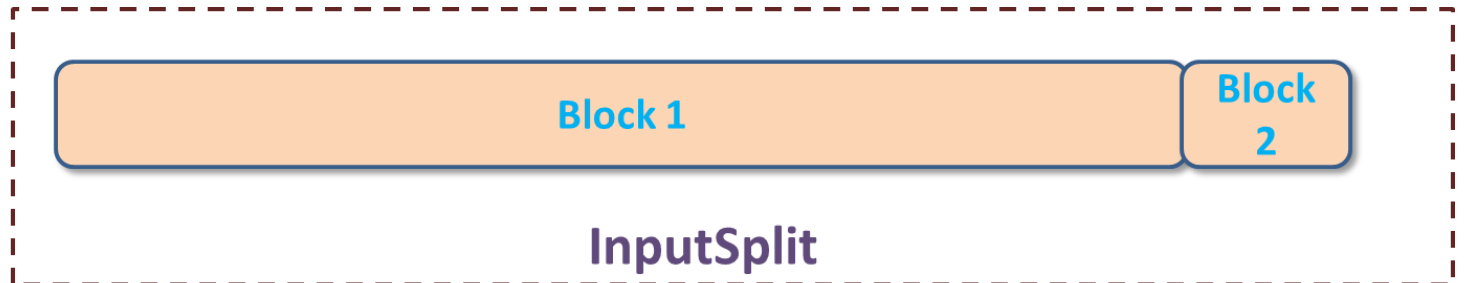
2 MB

Logical grouping
of blocks

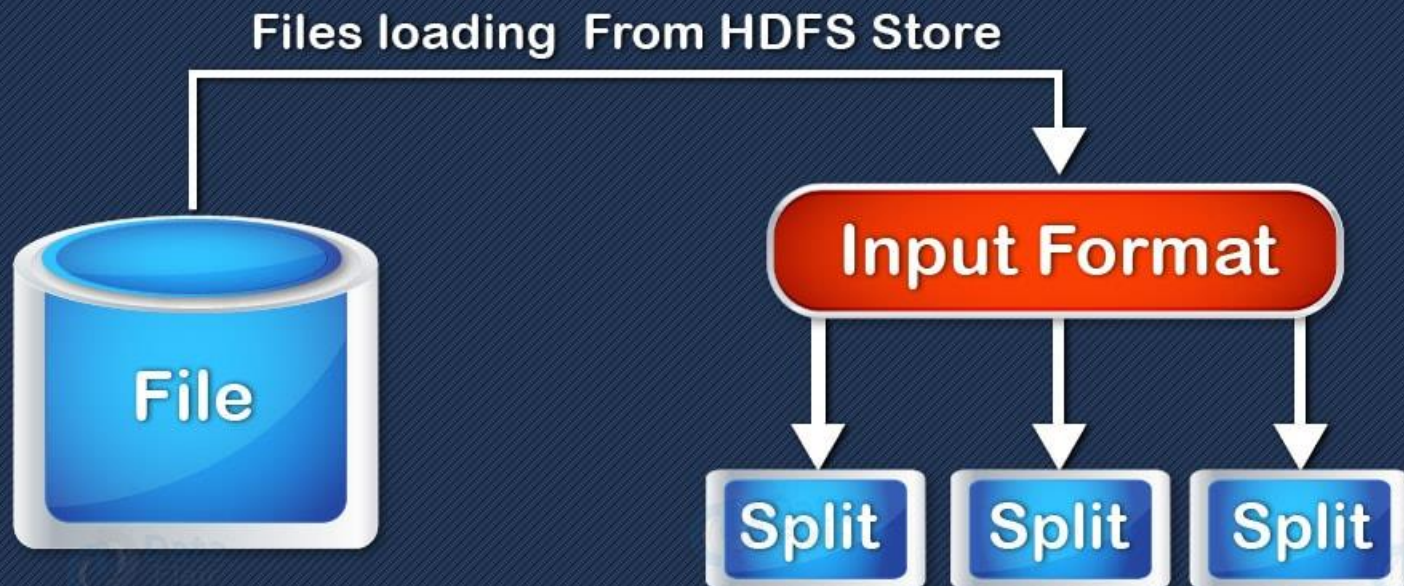
Block 1

Block
2

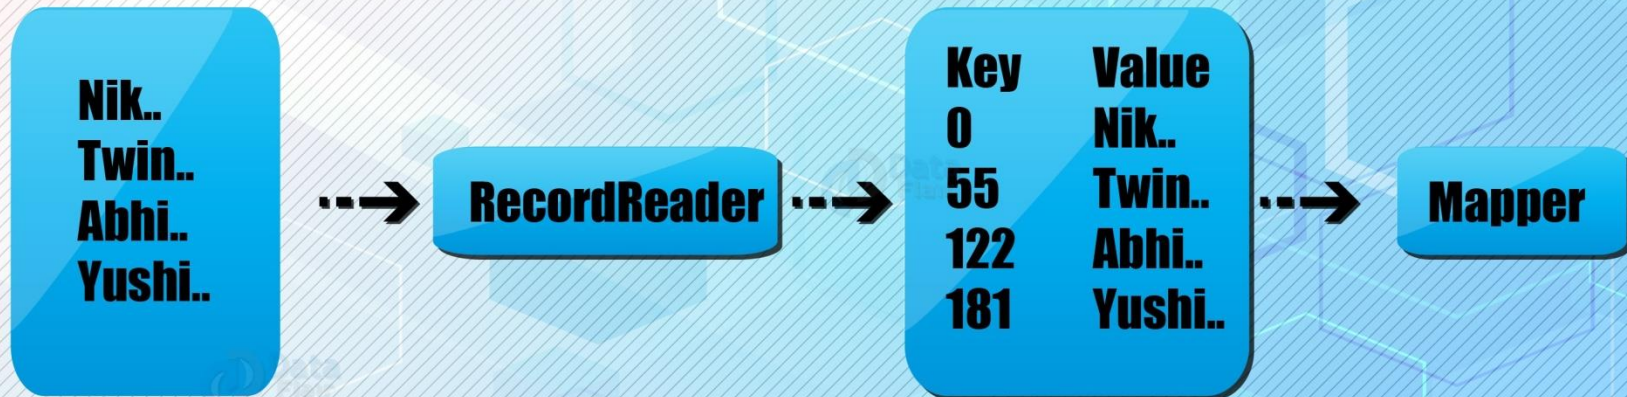
InputSplit



InputSplit in Hadoop MapReduce



RecordReader in Hadoop



It Converts Data into key Value Format

RecordReader

- Communicates with the **InputSplit** in Hadoop MapReduce and converts the data into key-value pairs suitable for reading by the mapper.
- By default, it uses TextInputFormat for converting data into a key-value pair.
- Communicates with the InputSplit until the file reading is not completed.
- It assigns byte offset (unique number) to each line present in the file. Further, these key-value pairs are sent to the mapper for further processing.

Process of conversion

- By calling '**getSplit ()**' the client calculates the splits for the job. Then it sent to the application master. It uses their storage locations to schedule map tasks that will process them on the cluster.
- After that map task passes the split to the **createRecordReader()** method. From that, it obtains RecordReader for the split.
- RecordReader generates record (key-value pair). Then it passes to the map function.

Mapper

- Processes each input record (from RecordReader) and generates new key-value pair, and this key-value pair generated by Mapper is completely different from the input pair.
- The output of Mapper is also known as intermediate output which is written to the local disk.
- The output of the Mapper is not stored on HDFS as this is temporary data and writing on HDFS will create unnecessary copies. Mappers output is passed to the combiner for further process.

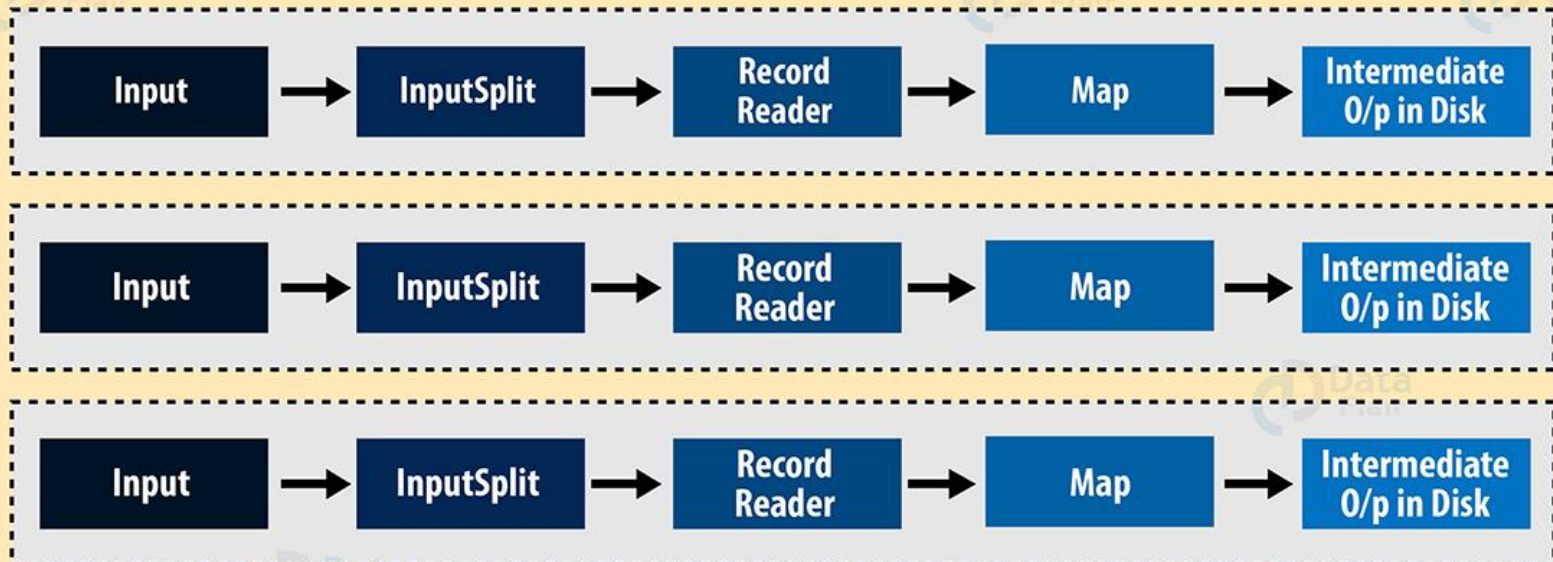
Disk Spill in Mapreduce

- Every map task is allocated some memory.
- Map output is stored in the buffer (default size 100MB but can be modified by `io.sort.mb` property).
- Disk Spill : When a certain threshold determined by `io.sort.spill.percent` (which is 80% by default) is reached the data is written on the local disk.
- And after the completion or failure of the job the temporary location used on the local filesystem get's cleared automatically. No manual clean up process is required, it's automatically handled by the framework.

Mapper



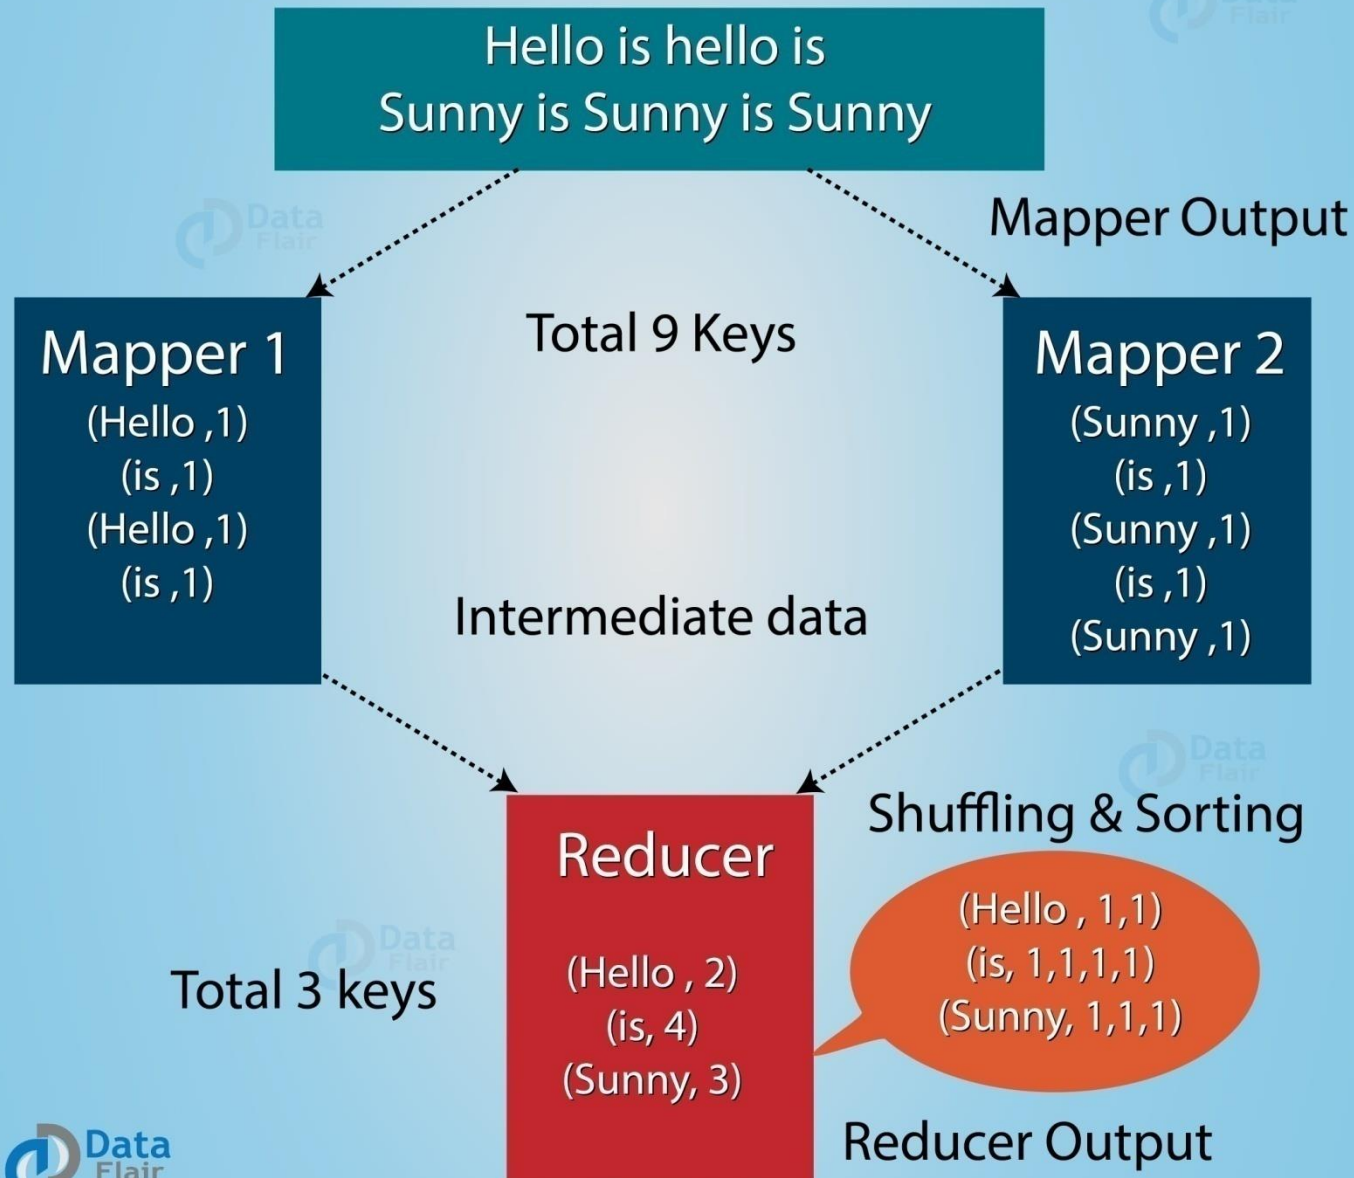
Mapper in Hadoop MapReduce



How many map tasks in Hadoop?

- The total number of blocks of the input files handles the number of map tasks in a program.
- The right level of parallelism is around 10-100 maps/node, although for CPU-light map tasks it has been set up to 300 maps. Since task setup takes some time, it's better if the maps take at least a minute to execute.
- For example, if we have a block size of 128 MB and we expect 10TB of input data, we will have 82,000 maps. Thus, the InputFormat determines the number of maps.
- Hence, **No. of Mapper = $\frac{\text{total data size}}{\text{input split size}}$**
- For example, if data size is 1 TB and InputSplit size is 100 MB then, No. of Mapper = $\frac{1000 \times 1000}{100} = 10,000$

MapReduce program without Combiner



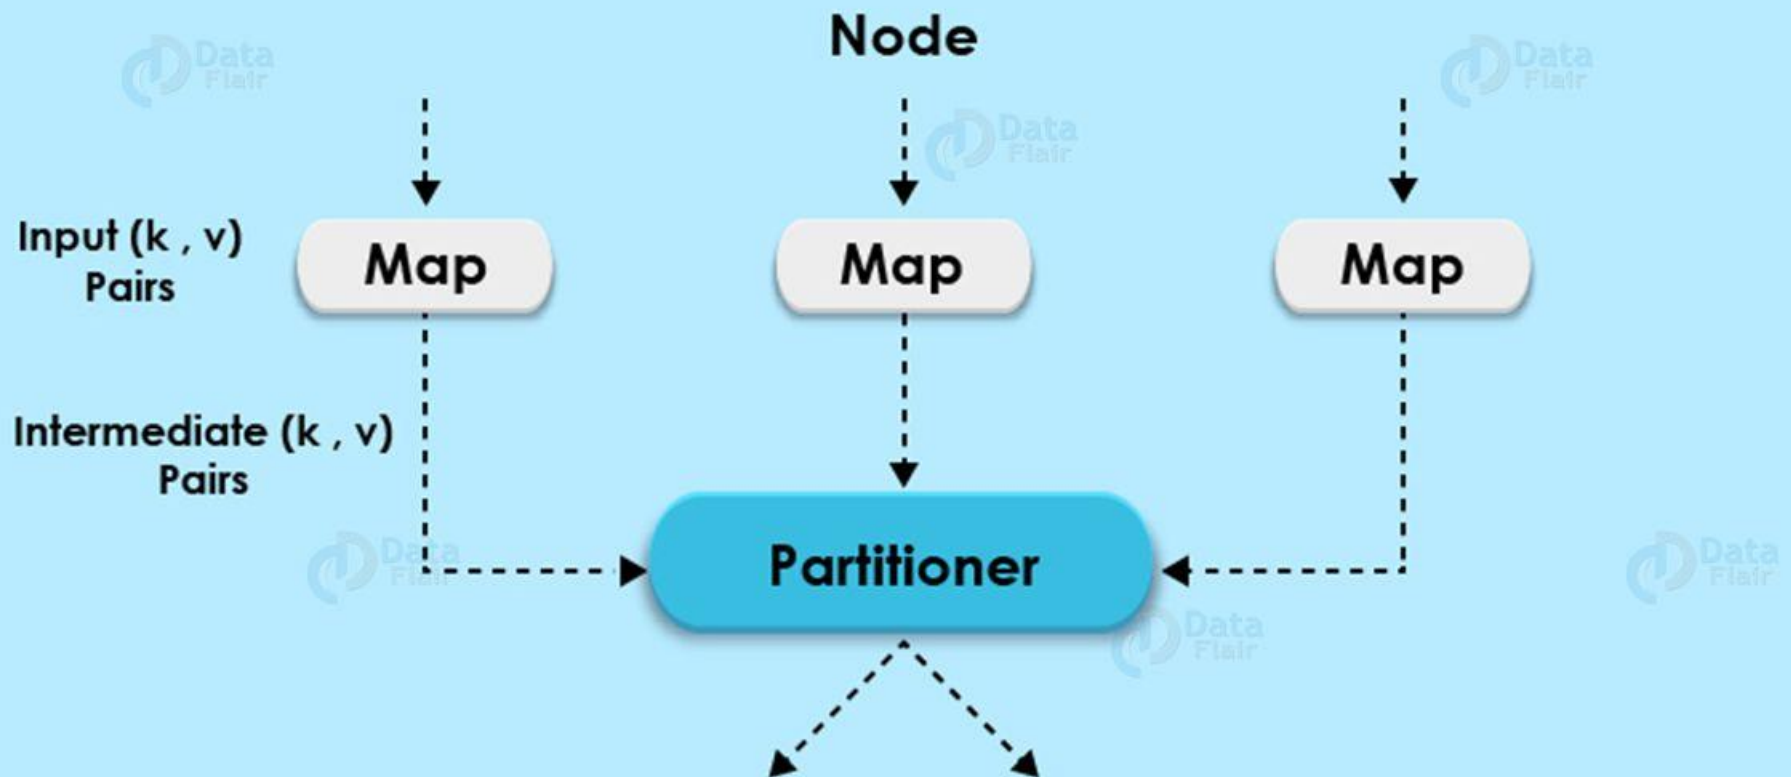
Combiner

- The combiner is also known as 'Mini-reducer'.
- Hadoop MapReduce Combiner performs local aggregation on the mappers' output, which helps to minimize the data transfer between mapper and **reducer**.
- Its usage is optional.
- Once the combiner functionality is executed, the output is then passed to the partitioner for further work.

Advantages of Combiner

- Use of combiner reduces the time taken for data transfer between mapper and reducer.
- Combiner improves the overall performance of the reducer.
- It decreases the amount of data that reducer has to process.

Partitioner in Hadoop MapReduce



Partitioner

- Comes into the picture if we are working on more than one reducer (for one reducer partitioner is not used).
- Takes the output from combiners and performs partitioning.
- Partitioning of output takes place on the basis of the key and then sorted. By hash function, key (or a subset of the key) is used to derive the partition.
- According to the key value in MapReduce, each combiner output is partitioned, and a record having the same key value goes into the same partition, and then each partition is sent to a reducer.

How many Partitioner in Hadoop?

- The total number of Partitioner depends on the number of reducers.
- Hadoop Partitioner divides the data according to the number of reducers.
- It is set by `JobConf.setNumReduceTasks()` method.
- The important thing to notice is that the framework creates partitioner only when there are many reducers.

Partitioning


- Partitioning allows even distribution of the map output over the reducer.
- e.g. $h(k) = k\%2$ divides data into two partitions
- Sometimes default partitioning may result in poor partitioning of data, if in data input in MapReduce job one key appears more than any other key.
- In such case, to send data to the partition we use two mechanisms which are as follows:
 - The key appearing more number of times will be sent to one partition.
 - All the other key will be sent to partitions on the basis of their `hashCode()`.
- If there is poor partitioning , the solution is to create Custom Partitioning.

Custom Partitioner Demonstration

Shuffling & Sorting in Hadoop

**Output
From
Mapper**

Ayush	432
Mona	467
Bittu	898
Disha	436
Disha	978
Ayush	345
Bretty	456
Mayank	967



Ayush	432
Ayush	345
Bittu	898
Bretty	456
Disha	436
Disha	978
Mona	467
Mayank	967

**Input
to
Reducer**

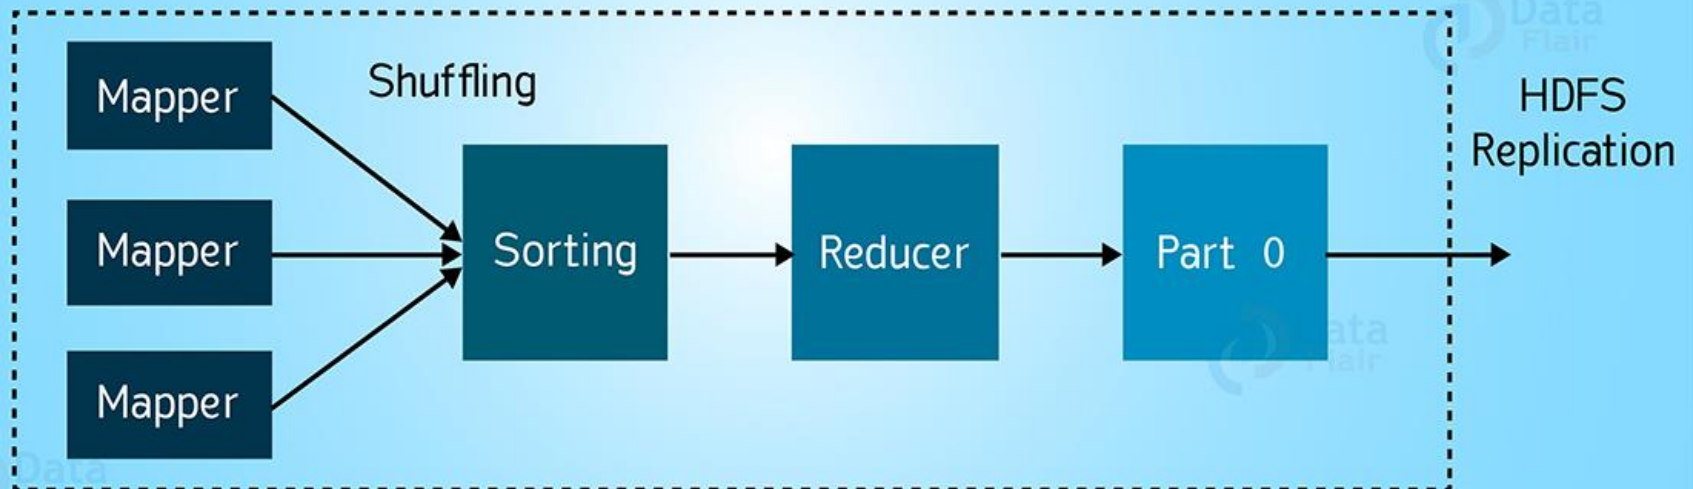
Shuffling and Sorting

- Shuffling and Sorting in Hadoop occurs simultaneously.
- The process of transferring data from the mappers to reducers is shuffling. It is also the process by which the system performs the sort. Then it transfers the map output to the reducer as input.
- MapReduce Framework automatically sort the keys generated by the mapper. Thus, before starting of reducer, all intermediate key-value pairs get sorted by key and not by value. Input from different mappers is again sorted based on the similar keys in different Mappers.
- If we want to sort reducer values, then we use a secondary sorting technique.

Reducer



Hadoop Reducer



Reducer

- It takes the set of intermediate key-value pairs produced by the mappers as the input and then runs a reducer function on each of them to generate the output.
- The output of the reducer is the final output, which is stored in HDFS.

Number of Reducers

- With the help of *Job.setNumreduceTasks(int)* the user set the number of reducers for the job. The right number of reducers are set by the formula, which is $0.95 \text{ or } 1.75 \times (\text{no. of nodes} \times \text{no. of the maximum container per node})$.
- With 0.95, all reducers immediately launch and start transferring map outputs as the maps finish. With 1.75, the first round of reducers is finished by the faster nodes and second wave of reducers is launched doing a much better job of load balancing.

Number of Reducers

- Assume, 100 reduce slots available in your cluster.
- With a load factor of 0.95 all the 95 reduce tasks will start at the same time, since there are enough reduce slots available for all the tasks.
 - This means that no tasks will be waiting in the queue, until one of the rest finishes.
 - Recommended when the reduce tasks are "small", i.e., finish relatively fast, or they all require the same time, more or less.
- With a load factor of 1.75, 100 reduce tasks will start at the same time, as many as the reduce slots available, and the 75 rest will be waiting in the queue, until a reduce slot becomes available.
 - Offers better load balancing, since if some tasks are "heavier" than others, i.e., require more time, then they will not be the bottleneck of the job, since the other reduce slots, instead of finishing their tasks and waiting, will now be executing the tasks in the queue.
 - Also lightens the load of each reduce task, since the data of the map output is spread to more tasks.

With the increase of the number of reducers

- Framework overhead increases.
- Load balancing increases.
- Cost of failures decreases.

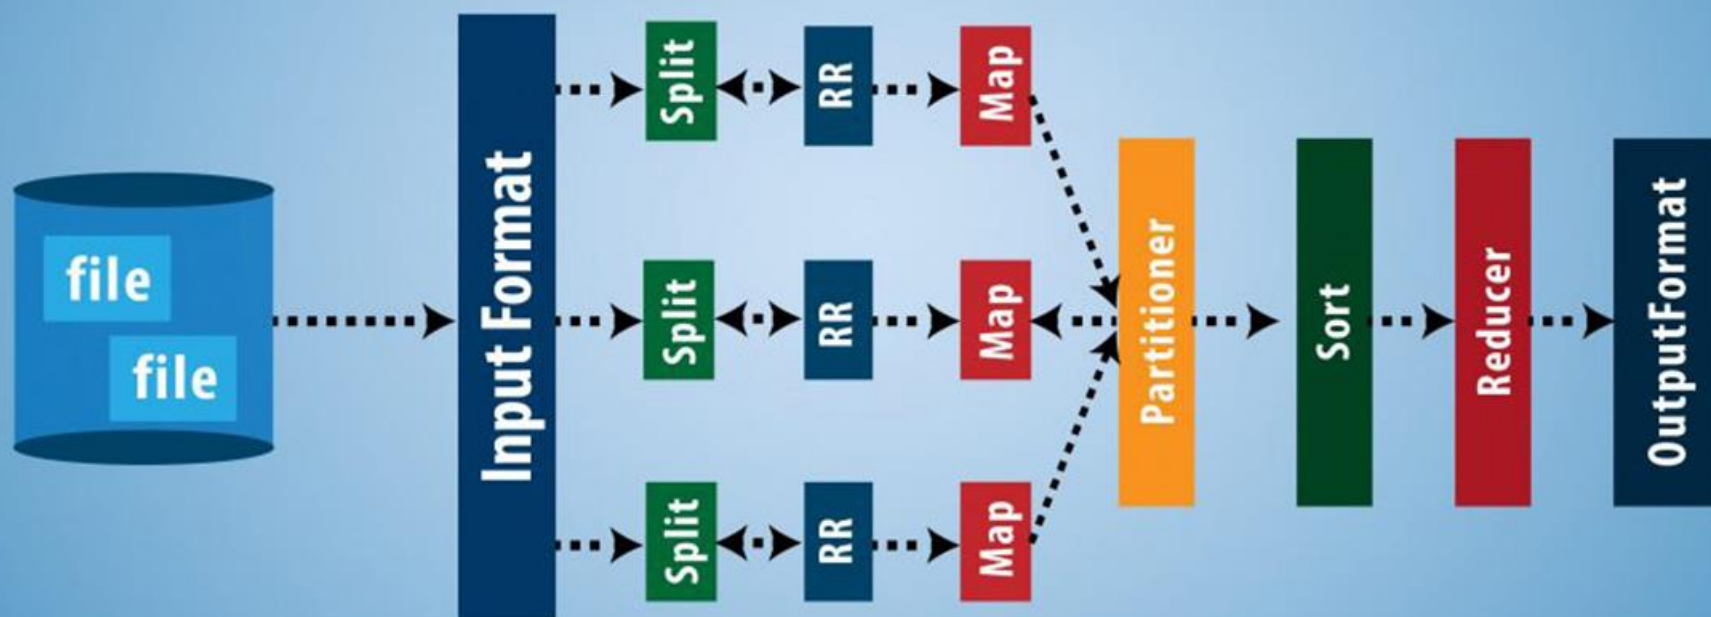
RecordWriter

- It writes these output key-value pairs from the Reducer phase to the output files.

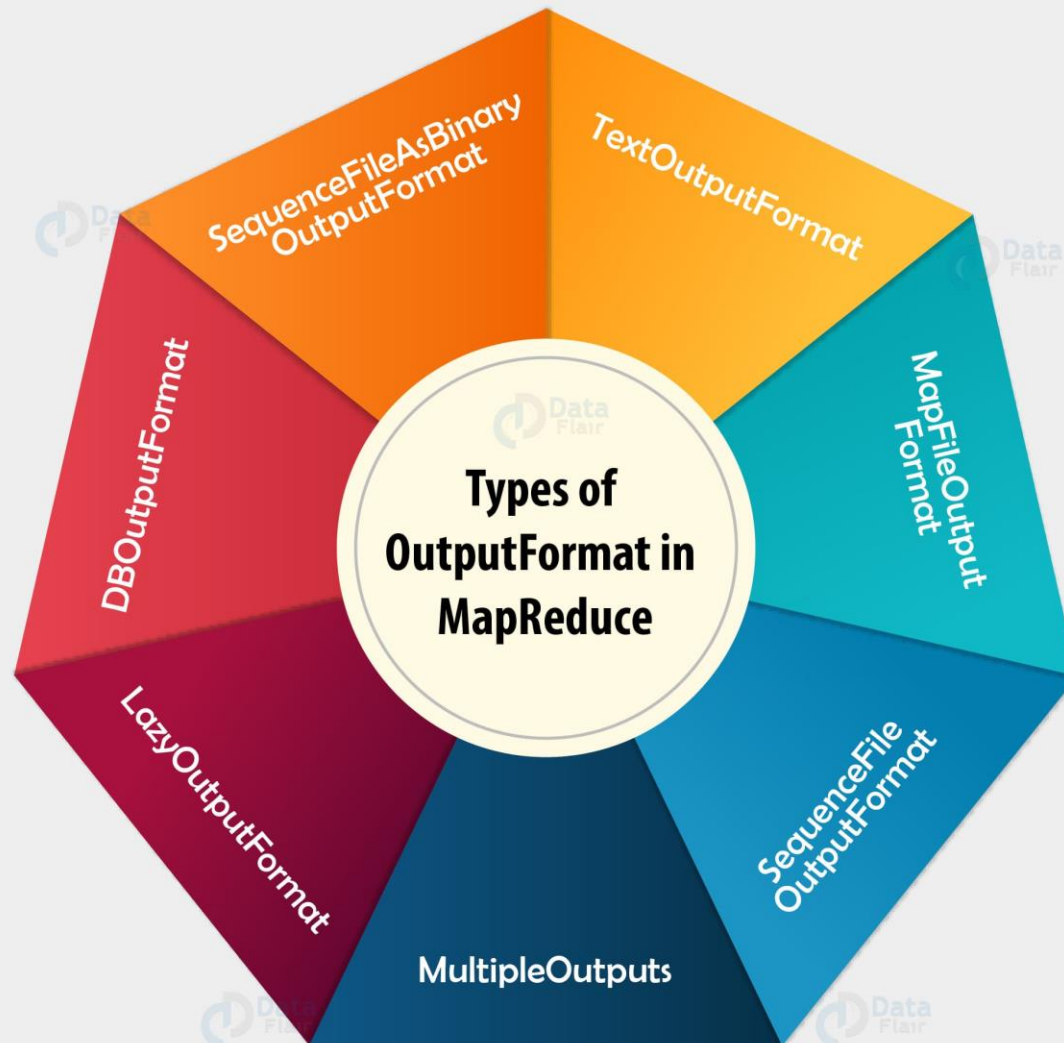
OutputFormat

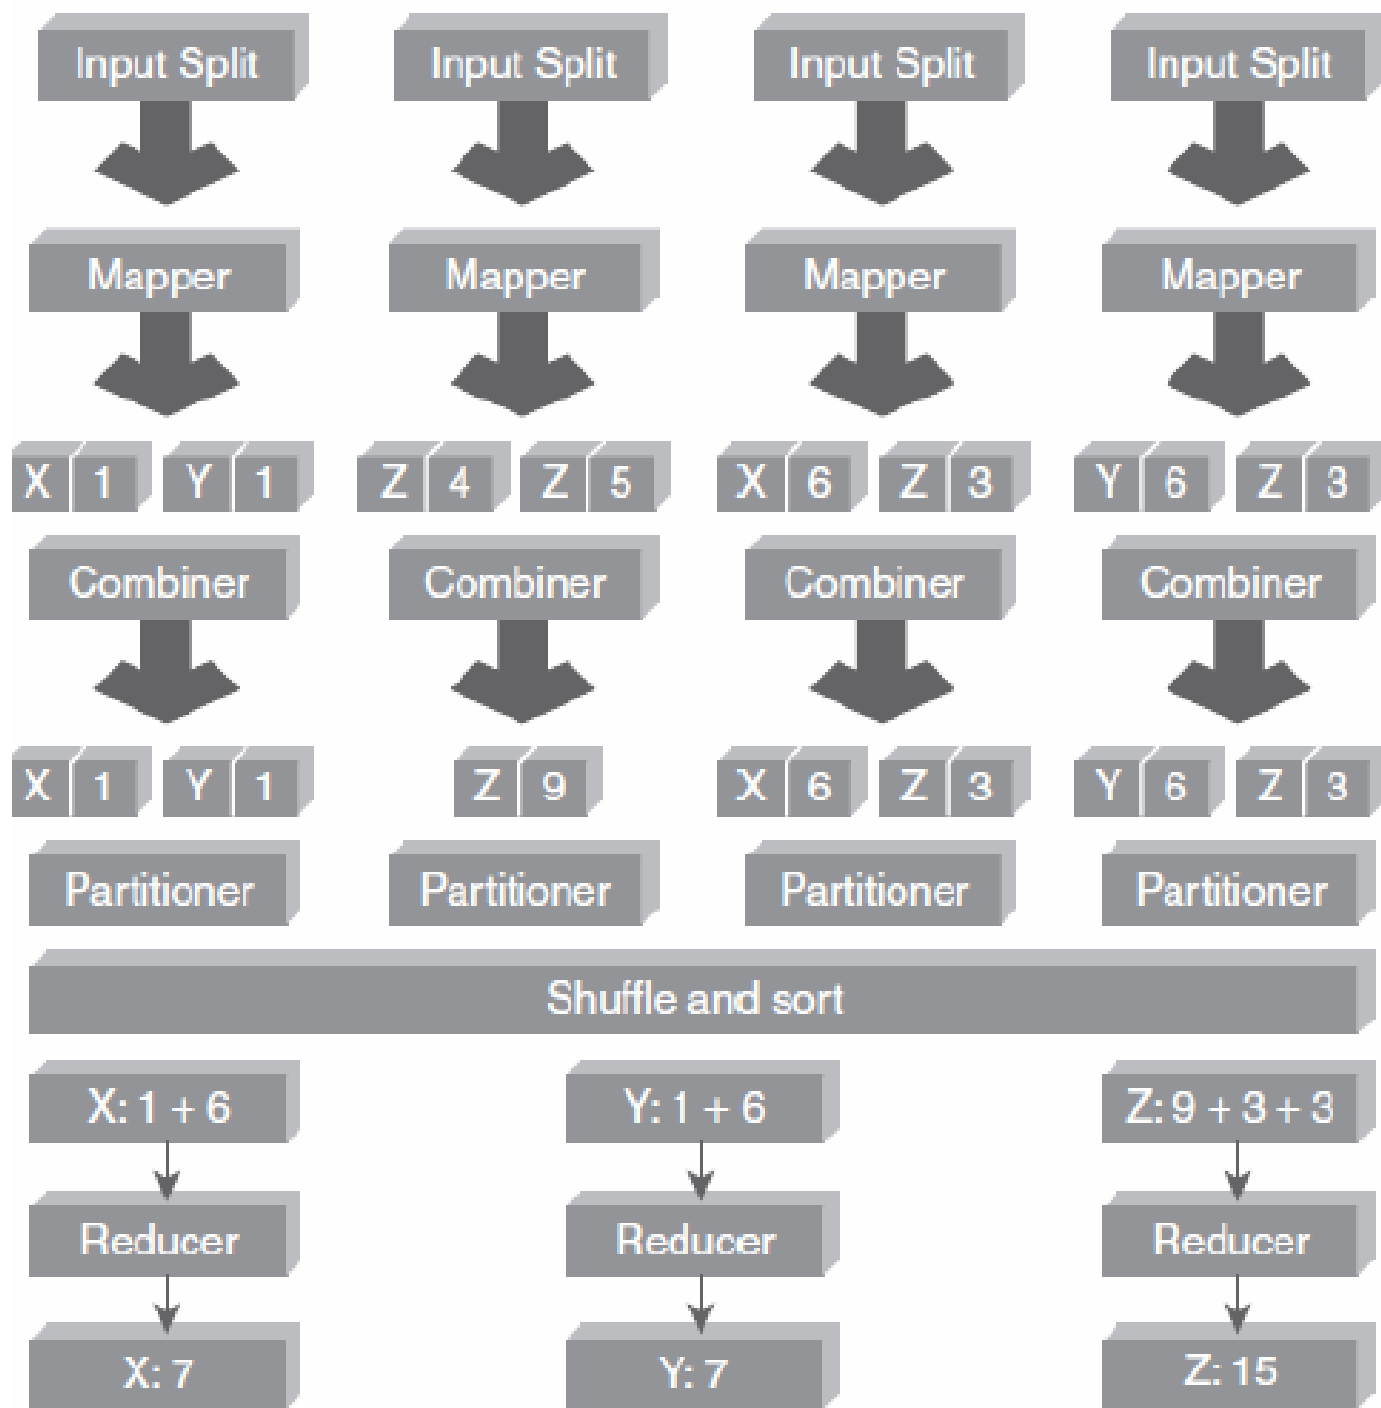
- The way these output key-value pairs are written in output files by RecordWriter is determined by the OutputFormat.
- OutputFormat instances provided by the Hadoop are used to write files in HDFS or on the local disk. Thus the final output of reducer is written on HDFS by OutputFormat instances.

OutputFormat in Hadoop MapReduce



Types of Hadoop Output Formats





References

- www.data-flair.training
- <https://techvidvan.com/tutorials/mapreduce-job-execution-flow/>
- <https://data-flair.training/forums/topic/in-map-reduce-why-map-write-output-to-local-disk-instead-of-hdfs/>
- <https://stackoverflow.com/questions/21980110/what-is-ideal-number-of-reducers-on-hadoop>

Thank you!